INKLOOM AI · ROUND 1 SUBMISSION · THE PLAN

PrismCheck AI

Intelligent pre-publication brand governance

Show the kit line, or stay silent.

Every finding cites an approved source | Every fix preserves creator intent | Every override teaches the brand kit

A review gate that sits between “someone made this” and “this goes live”, checking content against one approved Brand Kit and explaining every flag with a citation.

TEAM A2Z CREATORS

Venkata Narayana G V — Team Leader, Developer — venkatnarayana727@gmail.com
Mahesh Babu Ch — Research and Analysis — maheshmamtha242@gmail.com

3	4	100%	0
content formats reviewed	**cited checks per submission**	**of flags require an approved source**	**automatic overrides**
text · image · document	claims · tone · visual · access	citation or silence	human decision is final

Round 1 deliverable for the Inkloom AI brand-governance challenge. Prepared by Team A2Z Creators.

Contents

Eleven required sections, plus the supporting analysis that evidences each Round 1 judging criterion.

0	The one-page verdict	3
1	Understanding of the given problem statement	4
1b	Who else solves this, and where the gap actually is	6
2	The users affected by the problem	7
3	Proposed solution	9
3b	What the product actually looks like	11
4	The complete product workflow	12
4b	Three realistic catches	14
5	Main features we intend to build	15
6	Technology stack and technical plan	16
6b	Responsible design: privacy, security, oversight	18
6c	Risks, mitigations and a performance budget	20
7	Responsibilities of team members	21
8	Development plan across all three rounds	22
8b	The three improvements we name now	24
9	Expected impact	25
9b	Success metrics and evaluation plan	26
10	Future scope after the event	28
11	Complete team and participant details	29

Sections 1b, 3b, 4b, 6b, 6c, 8b and 9b are supporting analysis: competitor evidence, wireframes, worked examples, responsible design, risk, named improvements and success metrics.

0 The one-page verdict

Read this page if you read nothing else. The rest of the document is the evidence.

THE PROBLEM, IN ONE LINE

Organisations do not fail to govern their brand because they lack guidelines. They fail because the brand book is not the system of record at the moment content is published.

OUR ANSWER, IN ONE LINE

PrismCheck reviews text, images and documents against one approved Brand Kit before publication, explains every finding with a citation, and suggests fixes that keep the creator's intent.

WHY IT IS DIFFERENT

Other tools ask whether content feels off-brand. PrismCheck asks for the kit line: no source, no flag. And when humans repeatedly override a rule, it treats the rule, not the person, as the suspect.

How this document scores against the Round 1 rubric

25 pts · Understanding	Section 1 defines the failure modes line by line, names the exact gap we solve, and backs it with sourced evidence and a concrete organisational setting.
25 pts · Solution	Section 3 and 4 give a named, step-by-step walkthrough of one real submission, so a judge can see a person using the product, not just a feature list.
20 pts · Technical plan	Section 6 justifies every layer of the stack, names the hardest problem (an LLM that invents issues), and shows the architecture that prevents it.
15 pts · Originality	Section 3 states three mechanisms no other team is likely to have: Citation Lock, Intent Lock and Guideline Health.
15 pts · Feasibility	Section 8 commits to a deliberately small Round 2 build, deterministic checks first, with everything else explicitly deferred.

Sections follow the required submission order exactly: 1 Understanding · 2 Users · 3 Solution · 4 Workflow · 5 Features · 6 Technology · 7 Responsibilities · 8 Development plan · 9 Impact · 10 Future scope · 11 Team details.

1 Understanding of the given problem statement

What the brief actually asks, who feels it, and the specific slice we take responsibility for.

The brief is not a request for a brand checker. Growing organisations now publish through multiple employees, external agencies and AI tools, and none of them sit inside the central brand team. So the brief asks for something narrower and harder: a system that keeps a single approved source of truth in force at the exact moment content is published, by people and tools the brand team does not control.

Four failure modes named in the brief, and what each one costs

Incorrect claims and unsupported promotional promises

A creator or an AI tool invents a specification, price or guarantee that no approved source supports. The cost is legal and regulatory exposure, and the damage lands before anyone notices.

Outdated information and inconsistent language

Approved guidelines live in documents and folders that nobody opens at publish time, so last quarter's wording and last year's offer keep circulating.

Off-brand visuals and inconsistent asset use

Colour drift, unapproved logo handling and wrong typography are usually caught late, subjectively, or not at all, because visual checks are manual.

Accessibility failures

Low contrast and missing alternative text reach the audience, creating reputational and compliance risk that automated review could have caught deterministically.

Those failures map onto six detection areas the brief names explicitly: tone, terminology, messaging, visual consistency, accessibility and unsupported promotional claims. Each is a separate check in Section 5, and each must cite the kit item it tested against.

The gap we take responsibility for

One gate, one source of truth, one citation behind every finding.

We do not generate content, and we do not replace the brand team. We make the approved Brand Kit the only admissible evidence in the last review before publish.

Evidence that the problem is real

81%

of companies still deal with off-brand content

Lucidpress/Marq, 2019 · primary survey, n=200+

33%

upper-bound revenue lift from consistent branding

Lucidpress/Marq, 2019 · 23% in 2016

50%

of organisations created more content than the year before

Lucidpress/Marq, 2019 · primary

7–10 days

typical pre-governance approval cycle

Typeface, 2026 · practitioner guidance

How to read our evidence. Lucidpress/Marq (2019, n=200+) is primary research; Typeface, Aprimo and CrawlQ are labelled practitioner guidance, not research; WCAG 2.2 SC 1.4.3 and EU AI Act Article 50 are cited as specifications. We dropped stat-roundup figures we could not trace to a primary source: a plan is only as credible as its weakest citation.

The problem and the fix, in two pictures

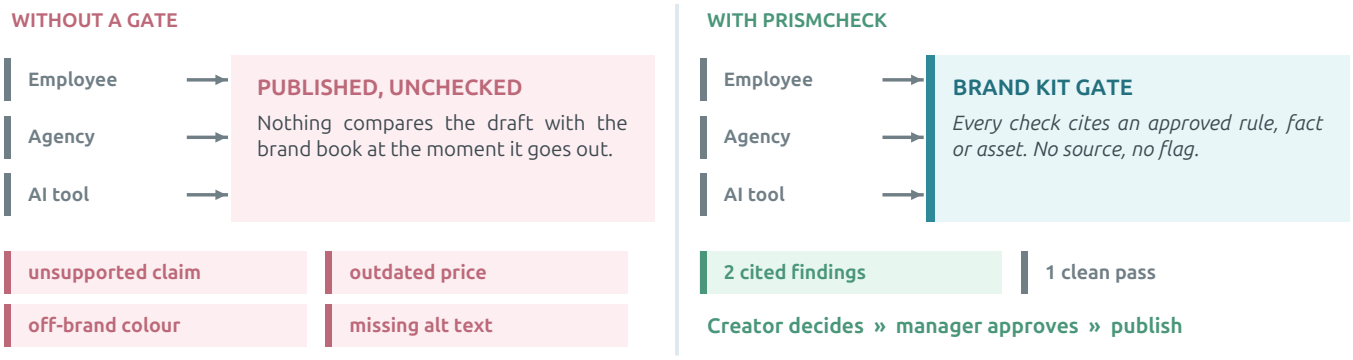


Figure 1 — Without a gate, the same three sources publish straight to the audience, and four kinds of error reach customers before anyone notices. With the gate, every draft is compared with the approved Brand Kit first, and nothing is flagged without a cited source.

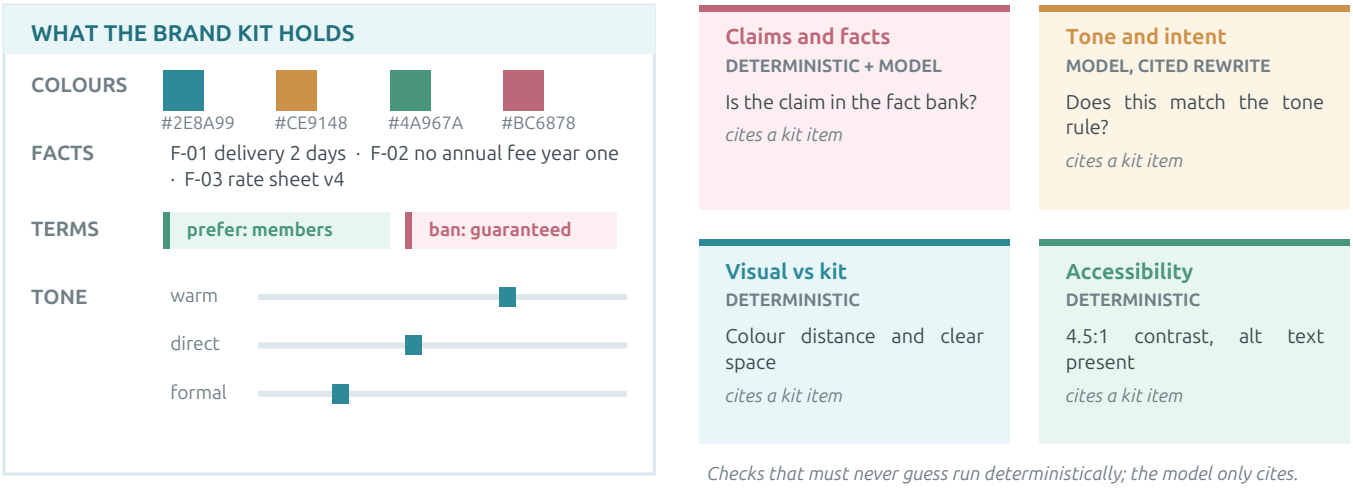


Figure 2 — The kit is small and structured on purpose: four kinds of approved material feed four checks. The checks that must never guess - terminology, colour distance and contrast - run deterministically, and the model only proposes a citation and a rewrite.

1b Who else solves this, and where the gap actually is

Judges see many brand-governance ideas. We checked the real market so our claim of originality is precise rather than loud.

Marq (Lucidpress)	Locked templates, brand asset management and approval workflows. Strong when content is created inside its templates. Gap for this brief: it does not audit a finished draft that arrived from somewhere else.
Adobe GenStudio	Brand compliance checks inside campaign production for ads and email. Strong inside Adobe workflows. Gap: it governs assets produced in its own pipeline, not drafts from employees, agencies or arbitrary AI tools.
Typeface	A brand hub plus AI generation with quality controls. Strong at governing what the model generates. Gap: PrismCheck governs what was produced, whoever produced it.
CrawlQ	Brand governance for AI content with multi-dimension scoring, knowledge-graph grounding and audit trails. Closest in philosophy. Gap: it scores generated output; PrismCheck refuses to flag without a citation, spans text, image and document, and treats human overrides as a signal about the guideline.
PrismCheck	The review gate: cite-or-silence findings against one approved Brand Kit, intent-preserving fixes, human approval, and override-driven guideline health. Designed for drafts of unknown origin.

We are not claiming the products above are poor. We are claiming they optimise a different moment in the workflow — creation, templating or generation — while the brief describes a review moment that sits after all of them. That is the space PrismCheck takes, and it is why our originality claim is about explainability rather than about features.

Why now

- Content volume is rising faster than review capacity: half of surveyed organisations produced more content year on year (Lucidpress/Marq, 2019).
- Generative tools removed the last natural checkpoint, because a draft can now be produced without a designer or writer who knew the rules.
- AI transparency expectations are becoming concrete: EU AI Act Article 50 obligations apply from 2 August 2026 for providers and deployers of certain AI systems, while in the US there is no general duty to disclose AI use but claim substantiation and endorsement disclosure rules still apply. A review trail is what makes either regime survivable.

2 The users affected by the problem

Three people feel this daily. Their needs drive the workflow in Section 4.

Content creator / copywriter

Need: Needs to ship fast without guessing.

Waits days for approval, gets subjective pushback with no citation, and learns the rule only after breaking it.

"I guess the tone and hope the manager doesn't catch it after the deadline."

Brand / compliance manager

Need: Needs to catch problems before they are live.

Spends review time on the same low-level errors, and usually sees the asset after it has already gone out.

"I keep seeing the same mistakes on things that go live before I review them."

Agency / freelance contributor

Need: Needs to know which source is current.

Works across several clients, each with a different brand document, and cannot tell which version is authoritative.

"Every client sends a different brand PDF, and I never know which one is current."

Persona quotes are illustrative composites for this plan; we will replace them with verbatim interview material during Round 2 user testing.

Organisational setting we design for

A mid-to-large organisation in D2C, fintech, edtech, SaaS, or campus and departmental communications. One central brand team, many distributed creators and agencies, and 30+ multi-channel assets a week across social, ads, landing pages and sales documents. Brand guidelines exist — as a PDF or a shared folder — but they are not in force at the point of publication.

Why this setting matters: it is big enough to need governance and small enough that one brand manager cannot manually review every asset, which is exactly the pressure the brief describes.

Research and supporting evidence we used

Desk research	Primary-source trade research published through 2026, including brand-consistency studies, trust research, governance guidance from platform vendors, and content-governance risk analyses. Every statistic in this plan carries its source.
Competitive review	We reviewed the incumbent approach — template lock-in, generation-side brand controls, and AI content scoring — to identify the gap PrismCheck occupies: an explainable, citation-backed review gate for drafts produced anywhere.
Practitioner observation	Persona pain points were built from publicly documented practitioner accounts of review friction, agency hand-offs, and late discovery of brand errors.
Planned primary research	In Round 2 we will run structured interviews with 3–5 creators and at least one brand reviewer, a short survey on where errors are discovered, and observation of one real approval cycle. Findings will be reported in the Round 2 submission.

3 Proposed solution

What PrismCheck is, the three mechanisms that make it different, and how each brief capability is covered.

PrismCheck AI is a pre-publication brand-governance platform. A creator submits a draft — text, an image, or a document. PrismCheck reviews it against the approved Brand Kit, returns findings that each cite a specific source, offers intent-preserving fixes, and routes the draft through a human approval workflow. It never generates content, never publishes, and never silently rewrites.

Three mechanisms that define the approach

1 Citation Lock

Cite, or stay silent

Every finding must carry a source_id pointing at a real guideline, fact or asset in the Brand Kit. A backend check drops any finding without a valid source before a user ever sees it. If the kit has no rule, PrismCheck says nothing instead of inventing an off-brand judgement.

2 Intent Lock

Fix the issue, keep the purpose

Before rewriting, PrismCheck extracts what the creator is trying to do — announce, reassure, warn, celebrate — then offers two alternative fixes that solve the flagged problem while preserving that intent. The creator always chooses; nothing is applied automatically.

3 Guideline Health

Overrides diagnose the kit

When a rule is repeatedly overridden or ignored across teams, PrismCheck surfaces it with an “ambiguous guideline” badge. Overrides become evidence that the guideline needs rewriting, not just proof that a creator got it wrong.

Every capability the brief asks for, and where PrismCheck provides it

Import or define guidelines, approved facts, terminology, reusable assets	Brand Kit builder: colour palettes with usage rules, approved facts with IDs, preferred and banned terminology, tone parameters, and reusable logo and asset entries.
Review text, images, documents before publication	One review flow with three parallel pipelines: text and caption, image and banner, and PDF/DOCX documents where text is extracted and embedded images are checked.
Detect tone, language, facts, claims, visual usage, accessibility	Four cited checks: claims and terminology, tone with intent analysis, visual consistency against the kit palette, and accessibility including contrast and alternative text.
Explain why flagged and identify the guideline or approved source	Each finding shows the issue, the guideline, fact or asset ID, the quoted source, severity, and the suggested fix.
Suggest corrections without overriding human decisions	Two intent-preserving rewrites per issue, with Accept fix or Keep mine and add context. Nothing is applied automatically and AI never publishes.
Comments, approvals, revision history, role-based review	Draft to Pending review to Approved or Changes requested, with threaded comments, revision history, and Creator, Manager/Reviewer and Admin roles.
Recurring problems and governance insights over time	Insights view with the most flagged rules and the most overridden rules, the latter carrying an ambiguous-guideline badge for brand team review.
Protect confidential materials and disclose AI use	Client-side masking of emails, phone numbers and internal tokens before any API call, confidential facts that never leave the organisation, and a plain-language AI disclosure in the interface.

VALUE PROPOSITION

PrismCheck is the last gate before publish — it tells a creator exactly which approved line their draft breaks, fixes it without erasing their purpose, and gives the brand team proof of which rules are failing them.

3b What the product actually looks like

A concept sketch of the two screens a judge will see first. Wireframe, not a screenshot — Round 2 replaces these with real captures.

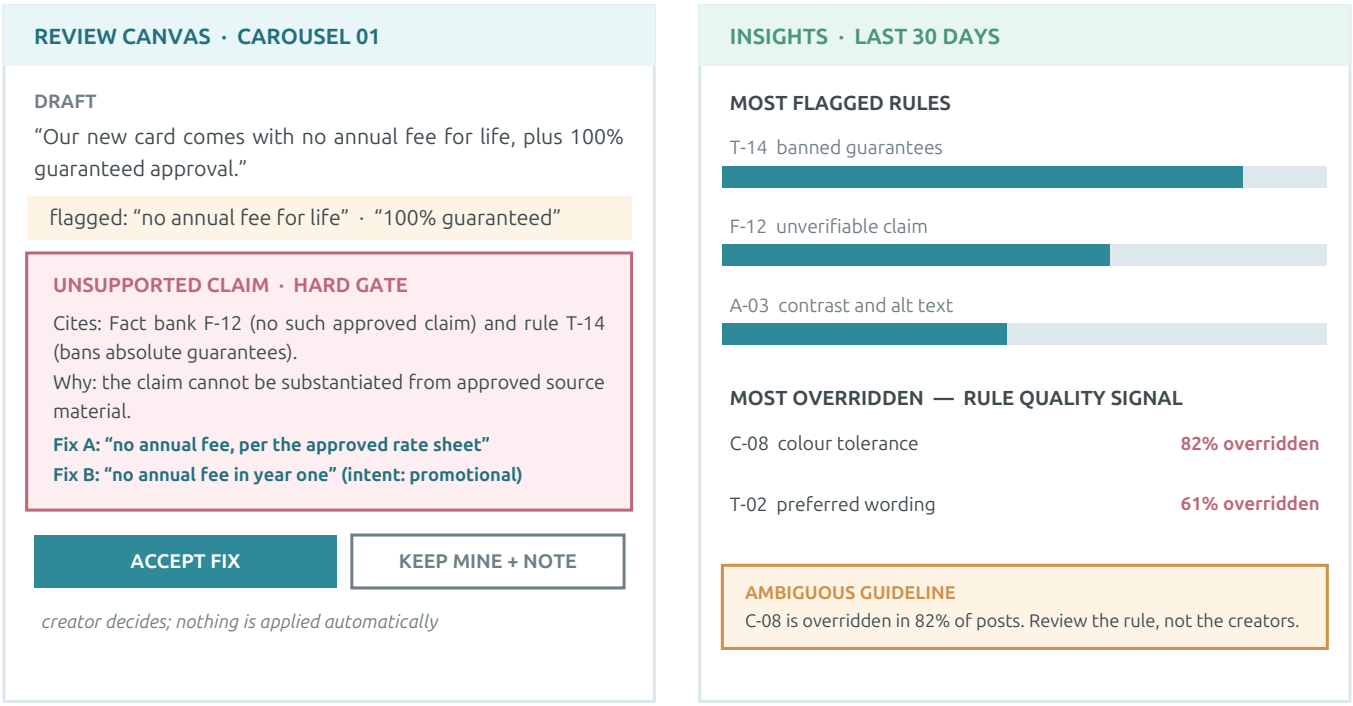


Figure 3 — Concept sketch. Left: every finding carries its citation and two intent-preserving fixes, and the creator holds the decision. Right: the same data drives governance — the most overridden rule is surfaced as a guideline problem, not a people problem.

Design decisions a judge can hold us to

- A finding never appears without a source chip. If a rule cannot be cited, the panel shows a "checked, no rule" state instead of a finding.
- Both actions are always available. There is no path where the system applies a change by itself.
- Silence is visible: the canvas reports how many rules were checked, so an empty screen is not mistaken for a broken one.
- Insights separates "rules broken" from "rules overridden", because those are different problems with different owners.

4 The complete product workflow

One diagram, then one walkthrough of a real submission from upload to insight.

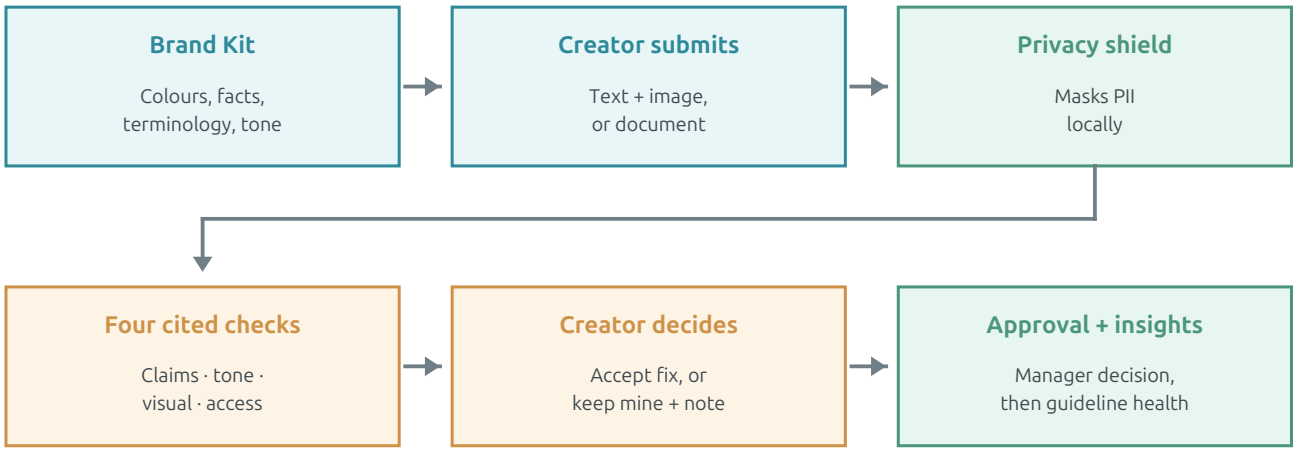


Figure 4 — PrismCheck review flow. The creator never loses control of the draft: the system only recommends, the creator decides, and the manager approves.

Walkthrough: a fintech creator submits a credit-card carousel

1 SUBMIT
A creator uploads a LinkedIn carousel announcing a new credit-card offer: caption copy plus a banner image.

2 PROTECT
Before anything leaves the browser, emails, phone numbers and internal tokens are masked by the privacy layer.

3 CHECK · CLAIMS
"No annual fee for life" is not in the approved fact bank. It is flagged as an unsupported claim, hard-gated, with the missing fact cited and two rewrites offered that keep the announcement intent.

4 CHECK · TONE
Intent is read as "urgent promotional offer". Two alternates preserve the urgency while removing the unsupported promise.

5 CHECK · VISUAL
Banner colours are compared to the kit palette in CIELAB space. One shade sits outside the approved range, so it is soft-flagged against the kit colour it drifted from.

Walkthrough continued — from the creator's decision to the governance insight

6 **CHECK · ACCESS**
Text over image fails the WCAG contrast ratio, and in-image text has no alternative text. Both are flagged with the measured ratio and a suggested description.

7 **DECIDE**
The creator accepts the claim rewrite and chooses “keep mine + add context” for the colour, recording the approved campaign exception as a note.

8 **APPROVE**
The draft moves to Pending review. The manager sees the fix, the exception note and the accessibility suggestions, then approves. Revision history records what changed and why.

9 **LEARN**
Insights shows this claim type flagged repeatedly this month, and the colour rule among the most overridden rules — badged as a possibly unclear guideline.

This is the whole product in one submission: a citation for every flag, a fix that keeps the creator's purpose, a human decision at every gate, and a governance signal at the end.

4b Three realistic catches

Concrete examples of what PrismCheck flags, what it cites, and what it suggests. These are the kinds of issues the brief describes.

Unsupported promotional claim

HARD GATE

DRAFT	Draft says: "Lasts 3x longer than the leading alternative."
CITED	Cited source: fact bank has no comparative specification and no such claim. The finding names the fact category that was checked and shows it is absent.
FIX	Suggested fixes (2): keep the benefit, drop the unverifiable multiple. e.g. "Built to last through everyday use" and "Longer wear between replacements, backed by our two-year warranty".
CONTROL	Creator chooses a fix or keeps their wording with a note. Manager must approve a hard-gated claim.

Banned promotional language

HARD GATE

DRAFT	Banner says: "100% guaranteed results — risk-free."
CITED	Cited source: terminology rule T-14 bans absolute guarantees and risk-free phrasing in this category. This check is deterministic string matching, not an LLM opinion, so it cannot be argued away or missed.
FIX	Suggested fixes (2): compliant alternatives that keep the confidence, e.g. "Backed by our results guarantee policy" and "Try it with our standard refund terms".
CONTROL	The manager sees the exact banned term and rule ID, so the discussion is about the rule, not about taste.

Accessibility failure

SOFT COACH

DRAFT	Banner has pale text on a pale photo, and the in-image headline has no description.
CITED	Cited source: accessibility rule A-03 requires a minimum contrast ratio for overlaid text and a description for text baked into images. Both are measured, not guessed: contrast is computed from luminance values.
FIX	Suggested fixes (2): darken the text or add a scrim to reach the threshold, plus a draft alt text generated from the OCR reading of the image text.
CONTROL	Soft guidance only: the creator decides, because some campaign art intentionally breaks contrast.

Note the pattern the rubric rewards: every example shows the issue, the exact approved source, the suggested correction, and who holds the final decision.

5 Main features we intend to build

Scoped deliberately. Everything else is explicitly listed as future scope in Section 10.

Authoritative Brand Kit	Structured entry of approved colour palettes with usage rules, approved facts with stable IDs, preferred and banned terminology, tone parameters, and reusable logo and asset records. This is the only source PrismCheck may cite.
Citation Lock engine	Structured findings that must carry a valid source_id. A middleware check discards any finding whose source is missing or does not match an ingested kit item, so an invented rule cannot reach the user.
Intent Lock rewriter	Extracts the creator's communicative intent, then produces two alternative corrections that resolve the flagged problem while preserving that intent.
Visual consistency inspector	Dominant colour extraction and CIELAB distance comparison against the kit palette, plus logo clear-space checking. Deterministic, local, and free of API cost.
Accessibility engine	WCAG 2.2 SC 1.4.3 requires a 4.5:1 contrast ratio for normal text and 3:1 for large text (18pt, or 14pt bold). PrismCheck computes the ratio from relative luminance, flags only values below the threshold, and drafts alternative text for text baked into artwork.
Collaborative review queue	Creator, Manager/Reviewer and Admin roles, threaded comments, approval states, and revision history with a readable diff of what changed and why.
Guideline Health insights	Two views: most flagged rules, and most overridden rules carrying an ambiguous-guideline badge that points the brand team back at the guideline itself.

Deliberately out of scope for Round 2: video and audio review, social platform publishing integrations, Canva and Figma extensions, multi-brand agency tenancy, and any custom model training. Naming them here is a scope decision, not an oversight.

6 Technology stack and technical plan

Why each layer exists, and the one hard problem this architecture is built to survive.

Frontend	React with Vite, Tailwind CSS and TanStack Table. Chosen for a fast side-by-side review canvas and dense, sortable findings and insights tables without a heavy framework.
Backend	FastAPI on Python 3.12. The review pipelines are naturally parallel — text and image checks run as separate async tasks — and Python is where the vision and NLP tooling lives.
Database	PostgreSQL via Supabase or self-hosted. Stores the kits, drafts, findings, approval state, comments and revision history, because the product needs structured audit data and role-scoped access, not blob storage.
AI layer	Google Gemini with structured JSON output enforced by schemas. Used for claim verification, tone and intent analysis, and rewrite suggestions — never as a free-form critic.
Fact retrieval	A private vector store (ChromaDB) holding only the organisation's approved facts. Retrieval happens inside the org boundary; the model may cite only what was retrieved.
Deterministic vision	Pillow and OpenCV for colour extraction, CIELAB distance and WCAG luminance maths. Local, free, reproducible, and independent of any model's opinion.
Auth and storage	Role-based JWT with Creator, Reviewer and Admin scopes. Uploads in a controlled local or S3-compatible store, never scattered through logs.

Section 6 continued — integrations, data sources, and the hard problem

Integrations and data sources

Data sources	Four only: the approved Brand Kit (colours, facts, terminology, tone, assets), the private approved-fact store used for claim retrieval, the uploaded draft itself (text, image or document), and the brand team's override history used for Guideline Health.
Integrations in Round 2	No external product integrations. The only outbound calls are to the disclosed AI provider for text analysis and alternative-text suggestions, plus the configured file store. This is a deliberate scope choice, not an omission.
Integrations deferred	Figma and Canva panels, social schedulers, CMS publishing webhooks and video pipelines are listed in Section 10 as post-event scope.

Data flow in one line

Kit and draft in » client-side masking » parallel text, image and document pipelines » Citation Lock validation » cited diff in the review canvas » approval state machine » insights and guideline health.

The hardest problem, stated plainly

THE RISK

Ask a general model to “check this for brand compliance” and it will confidently invent rules that are not in the kit, then justify them in fluent prose. That is the exact opposite of the explainable, reliable results the brief demands, and it is the failure mode a judge is most likely to probe.

THE ARCHITECTURE THAT ANSWERS IT — FOUR LAYERS OF DEFENCE

1. Only the kit is injected as context, so the model has nothing else to cite.
2. Every finding is a validated object requiring a source_id that matches an ingested kit item.
3. A middleware gate drops any finding with a missing or unmatched source before the user sees it.
4. Safety-critical checks never touch the model: banned terms use trie or regex matching, and colour and contrast use deterministic maths on the image itself.

The result is a system that under-reports rather than fabricates. When PrismCheck has no rule to cite, it stays silent and surfaces the ambiguity through Guideline Health instead of guessing.

6b Responsible design: privacy, security, oversight

The brief asks for this explicitly, so it gets its own page rather than a footnote.

Privacy

Emails, phone numbers and internal tokens are masked in the browser, before any content leaves the creator's machine. Brand Kit facts marked confidential never travel outside the organisation. No draft, image or finding is retained for model training, and the UI states this in plain language rather than in a policy page.

Security

Role-based access with three scopes — Creator, Reviewer/Manager, Admin — enforced at the API, not hidden in the interface. Kits, drafts, findings, comments and revision history live in PostgreSQL with role-scoped reads. Uploads sit in a controlled store with no content in application logs. Every approval action is attributable to a named user.

Human oversight

The product balances automation with human control by splitting checks in two. Hard gates — unsupported claims, contradicted facts, banned legal phrasing — cannot be cleared by the creator; they require an authorised manager. Soft coaching — tone, colour drift, accessibility — is advice the creator may accept or decline with a recorded reason. AI never publishes, never rewrites silently, and never overrides a human decision.

AI and external service disclosure

SHOWN IN THE PRODUCT INTERFACE

"Text analysis and alternative-text suggestions use Gemini. Colour and contrast checks run locally on your machine. No content is stored as training data. Brand kit facts marked confidential stay in this organisation."

Where the law actually stands, stated carefully. In the EU, Article 50 of the AI Act sets transparency obligations for providers and deployers of certain AI systems and applies from 2 August 2026. In the US there is no general federal duty to disclose AI-generated content, but advertising claim substantiation and endorsement disclosure rules still apply to whatever is published. So PrismCheck does two different things and does not blur them: it discloses its own AI use in the interface as a trust measure, and it treats claim substantiation as a hard gate, because that is where the legal exposure actually sits.

Section 6b continued — the brief's expected outcome, clause by clause

What the brief expects, and where this plan delivers it

A functional product	Round 2 ships a working review flow end to end: upload, four cited checks, creator decision, manager approval, revision history, insights. Round 3 hardens it.
Approved source material	Nothing is flagged unless it resolves to an approved guideline, fact or asset in the Brand Kit. The Citation Lock enforces this in code, not in prompt wording.
Understandable findings	Each finding shows the issue, the source ID, the quoted rule, a severity and a suggested fix — readable by a non-technical creator without a glossary.
Human review	Every draft passes through a named reviewer before it can be marked ready. Hard-gated findings cannot be self-cleared by the creator.
Realistic workflows	Three real formats (caption, banner image, PDF or DOCX one-pager), a worked carousel example, and approval states that mirror how brand teams already work.

6c Risks, mitigations and a performance budget

What could go wrong, what we do about it, and the numbers we will be measured against.

The model invents an issue	Highest risk in this category. Mitigation: Citation Lock — unsourced findings are dropped by middleware before display; terminology and colour checks are deterministic, not model-driven.
The kit is too thin to review	The system stays silent and records the gap. Mitigation: an onboarding checklist and a “checked, no rule” state so silence is visible rather than mysterious.
Accessibility false positives	Brand art sometimes breaks contrast on purpose. Mitigation: accessibility is soft coaching only, and an override records the reason for the next reviewer.
Colour drift false positives	Campaign palettes legitimately differ. Mitigation: a tolerance band plus per-campaign approved exceptions stored in the kit rather than argued in chat.
Document parsing fails	A corrupted or scanned PDF should not pass silently. Mitigation: an explicit unreadable state that blocks approval and asks for a text version.
Confidential material leaks	Mitigation: masking happens client-side before any call; confidential facts are excluded from model context; no content is retained for training.
Scope overrun in the build window	Mitigation: deterministic checks ship first, the model path second, workflow third, insights last — with a fixed cut list if time runs short.

Performance and cost budget, stated as targets

We are not claiming measured numbers in Round 1. These are the targets we will validate and report in Round 2, with the measurement method named so the claim is checkable.

Deterministic checks	Under 500 ms per image on a laptop CPU. Measured locally with timing instrumentation around the OpenCV and contrast functions.
Text review with model	Target p50 under 5 s for a 300-word draft, hard timeout at 20 s returning partial deterministic results rather than an error page.
Model calls per review	Budget of three: claim verification, tone and intent, rewrite generation. Measured by counting outbound calls per submission.
Cost per review	Bounded by the three-call budget and by trimming kit context to the retrieved facts, which keeps token cost predictable at low volume.

7 Responsibilities of team members

Two members, five workstreams, with clear ownership so any judge can ask either of us to demonstrate our part.

Venkata Narayana G V

Team Leader · Developer

Owns the product end to end: FastAPI architecture, PostgreSQL schemas, the approval state machine and role model, the Citation Lock middleware, the review canvas, and the final demo and narrative.

venkatnarayana727@gmail.com

Mahesh Babu Ch

Research and Analysis

Owns research and analysis: user research and persona validation, the evidence pack and source list, competitive and gap analysis, findings-quality analysis including false-positive measurement, and validation of accessibility thresholds against WCAG.

maheshmamtha242@gmail.com

Workstream ownership

Backend and workflow	PrismCheck API, database schema, approval state machine and revision history. Owner: Venkata Narayana.
AI and citation contract	Gemini structured output, fact retrieval, intent extraction and rewrite quality, source validation. Owner: Venkata Narayana.
Vision and accessibility	Colour distance and contrast maths, in-image text reading and alternative-text drafts. Implemented by Venkata Narayana; thresholds validated by Mahesh Babu.
Research and analysis	User interviews, persona validation, evidence sourcing, gap analysis, and findings-quality testing. Owner: Mahesh Babu.
Interface and demo	Review canvas, kit builder, insights views, and the demonstration script. Shared, led by Venkata Narayana.

During Round 2 user testing we will record which features each member shipped, so the contribution claim stays checkable against the repository history.

Why this team can build it

- Skill coverage matches the build: one member owns backend, data model, citation middleware and interface; the other owns research, evidence, findings-quality analysis and accessibility threshold validation. No required workstream is unowned.
- We build the smallest thing that works first. Deterministic checks ship before any model call, so a demonstrable product exists early even if the AI layer needs rework.
- Both members can demonstrate their own part to a judge without notes, which is the standard Round 2 and Round 3 judging puts us to.
- We write down what we cut and why. A smaller product that runs end to end scores better on this rubric than an ambitious one that half-works.

8 Development plan across all three rounds

A focused build order: the deterministic backbone first, then the cited AI layer, then the workflow and insights.

ROUND 1 **The plan** *Complete*

- Define the Brand Kit schema and the finding model with a mandatory source_id.
- Specify the approval state machine and the three roles.
- Specify the privacy layer and the deterministic checks.
- Publish the evidence pack, walkthrough and examples.

ROUND 2 **The build** *Core milestone*

- Deterministic checks first: banned terms, colour distance, contrast, logo clear space.
- Citation Lock: structured output, kit-scoped context, middleware drop of unsourced findings.
- Fact retrieval over the approved fact bank, then intent extraction and rewrites.
- Approval workflow, revision history, and privacy plus AI-disclosure UI.
- Usability testing with 3–5 creators and one reviewer.

ROUND 3 **The finish** *Polish*

- Harden Guideline Health and the ambiguous-guideline badges.
- Improve rewrite quality and intent classification from real test feedback.
- Reliability pass: error states, empty kits, missing sources, malformed files.
- Demo video, documentation, known limitations and the post-event roadmap.

Build-order logic, and why it is feasible

- The deterministic checks are built first because they work without any model call. If the AI layer is delayed, the product still reviews terminology, colour and contrast.
- The Citation Lock is built second because it is the backbone of explainability. If it works, every later feature composes around it.
- Workflow and insights come last because they are ordinary application logic once findings are trustworthy.
- The scope is fixed at one Brand Kit, one submit flow, three formats, four checks, one approval workflow and one insights page.

Section 8 continued — how we will spend the build window

Allocation of the build window

First 25% of the window	Skeleton that runs end to end: upload a draft, store it, show a findings list, mark it pending. Ugly is acceptable; unusable is not.
25% to 60%	The deterministic checks, the Brand Kit builder, and the Citation Lock with source validation. This is the load-bearing work.
60% to 85%	Model path for claims, tone and intent, then the approval workflow, revision history and the insights page.
Final 15%	Hardening, labelled-set evaluation, usability testing with 3–5 creators, then the demo recording and documentation.

If the window is shorter than assumed, the cut order is fixed in advance: insights first, then intent rewrites, then document support. The cite-or-silence review gate is never cut, because without it there is no product.

8b The three improvements we name now

Round 2 scores "the three improvements you named" (15 points) and Round 3 scores "the three improvements delivered" (20 points). So we name them in the plan, each with a reason, a change and a measurement.

1 Exception memory

Stop re-arguing exceptions that were already approved

WHY A creator overrides the same rule every campaign, so the flag becomes noise and the reviewer spends time re-reading the same note. Approved overrides are currently remembered only in a comment thread.

CHANGE An approved override can be promoted to a scoped exception rule with an owner and an expiry date. Future findings cite the exception instead of re-flagging the same issue, and the exception is visible in Guideline Health so it can be reviewed or retired.

MEASURED BY Repeat-override rate on the same rule, and median notes per asset before approval.

2 Inline evidence panel

Show the rule text, not just its identifier

WHY A citation ID is precise but not readable at speed. A creator who is not a brand specialist still has to ask what rule T-14 actually says, which recreates the conversation the product was meant to remove.

CHANGE The finding panel expands the cited kit item inline, with the matching phrase highlighted and the source version shown. One click from the flag to the exact approved line, with no separate document to open.

MEASURED BY Time from finding shown to creator decision, plus a comprehension rating in usability testing.

3 Batch campaign review

One decision for a campaign, not twelve separate ones

WHY Campaigns ship eight to twelve variants. Reviewing them one at a time rebuilds exactly the bottleneck the brief describes, even if each individual review is fast.

CHANGE Select several assets, run the four cited checks across all of them, then approve or request changes per asset or as one group decision that writes a single audit entry covering the set.

MEASURED BY Reviews completed per hour, and reviewer-reported effort per campaign.

If one of these slips, we will say so in Round 2 rather than renaming a shipped feature as an improvement. The three are ordered by expected impact, so the list is also the cut order.

9 Expected impact

Grounded claims, tied to the evidence in Section 1. No invented percentages.

Faster routine review

Governance research reports pre-governance cycles of 5–7 revision rounds and 7–10 days, improving to 2–3 rounds and 2–4 days once governance is in place. PrismCheck targets that shift by removing the basic errors before a human sees the draft.

Fewer unsupported claims in market

Hard gates stop unverified specifications, invented comparisons and banned guarantees from being marked ready without an authorised manager decision, which is the failure the brief opens with.

A brand kit that improves itself

Guideline Health converts override patterns into a maintenance backlog for the brand team. Instead of the same rule being broken forever, an unclear rule gets rewritten once and stops generating findings.

How we will measure it during Round 2 and 3

Findings quality	Percentage of findings a reviewer agrees with, measured on a labelled set of 30 real drafts.
False positive rate	Share of findings rejected as wrong. PrismCheck is designed to under-report rather than fabricate, so this is the number that proves the Citation Lock works.
Time to first review	Difference between submission and the first substantive reviewer action, compared against the current manual flow.
Citation coverage	Percentage of findings carrying a resolvable source ID. Target: 100%, enforced by the middleware.

9b Success metrics and evaluation plan

Impact claims are only credible if the measurement is designed before the feature is built.

Citation coverage	100% of displayed findings resolve to a live kit item. Enforced in code, so this is a correctness gate, not a target.
Reviewer agreement	Target 80% or higher: the share of findings an independent reviewer marks as valid, on a labelled set of 30 real drafts.
False positive rate	Target 10% or lower. PrismCheck is deliberately built to under-report, so this number is the proof the Citation Lock works.
Silence accuracy	Target 90% or higher: drafts with genuinely no rule violation produce no findings. This is the test that catches over-flagging.
Time to first review	Target a 50% reduction against the team's current manual baseline, measured as submission-to-first-substantive-comment.
Guideline actions	At least three rules rewritten per quarter from Guideline Health signals — the measure that the product improves the kit, not just the draft.

Evaluation protocol

- Build a labelled set of 30 real drafts at three quality levels: clean, mildly off-brand, and contains an unsupported claim.
- Two independent raters mark each finding as valid or invalid; disagreements are discussed and the rubric tightened, not averaged away.
- Report inter-rater agreement (Cohen’s kappa, target 0.6 or above) alongside the accuracy figures, so the numbers are interpretable.
- Run the same set against a naive “ask the model to check brand compliance” baseline, and publish both results. If our approach does not beat the baseline, that is the finding.
- Test with 3–5 creators and at least one reviewer, and report the friction they hit rather than only the tasks they completed.

Section 9b continued — the demo we will give a judging panel

The three-minute demo we will give

0:00	Open a pre-loaded Brand Kit: four colours, nine facts, one banned term, one tone definition.
0:30	Upload a carousel whose caption contains an unverifiable claim and whose image uses an off-palette shade with low-contrast text.
1:00	The finding appears with its citation, the quoted rule, and two rewrites that keep the promotional intent.
1:30	Accept one fix, keep the colour with a note; the draft moves to Pending review.
2:00	Switch to the reviewer account and approve; revision history shows exactly what changed and who decided.
2:30	Open Insights and show the overridden colour rule badged as a possible ambiguous guideline.
2:50	Close on the empty state: a clean draft produces no findings, and the canvas states how many rules were checked.

The last step matters most. A governance tool that only ever shows problems teaches judges nothing about whether it can stay quiet — and staying quiet is the entire thesis.

10 Future scope after the event

Deliberately deferred. Nothing here is promised for the build window.

Creative-suite integrations

Figma and Canva panels that lint inside the tool where designers actually work, instead of asking them to upload a finished asset.

Video and audio review

Frame-aware logo tracking, on-screen text checks, and transcript analysis for reels and short-form video.

Verified publishing hand-off

One-click publishing to social schedulers and CMS platforms once an asset is approved, carrying the approval record with it.

Deeper compliance profiles

Sector-specific rule packs for regulated industries such as financial services and healthcare, where claim wording carries legal weight.

If the build window allows only part of one of these, we will ship the deterministic half and document the rest. Scope honesty scores better than an unbuildable wish list.

11 Complete team and participant details

Names exactly as they should appear on event records and certificates.

TEAM NAME

A2Z Creators

Member 1	Venkata Narayana G V
Role	Team Leader, Developer
Email	venkatnarayana727@gmail.com
Member 2	Mahesh Babu Ch
Role	Research and Analysis
Email	maheshmamtha242@gmail.com

SUBMISSION NOTES

This PDF is the Round 1 plan for the Inkloom AI brand-governance challenge. All work described is planned for the official build window. Team names and spellings above are final for certificate purposes; please do not adjust them after submission.

Declarations

- The product will be built during the official build window. Research, planning and user interviews were conducted beforehand, as permitted.
- External services used at build time will be disclosed in Round 2: Google Gemini for text analysis and alternative-text suggestions, a private vector store for approved facts, and a local or S3-compatible file store.
- Brand kit facts marked confidential stay inside the organisation; no customer content is used as training data.
- AI assistance disclosure: this plan document was prepared with AI writing and layout assistance, which we state openly. The product itself will be built during the official build window. Significant AI coding assistance during that window, and any reused template or boilerplate code, will be disclosed in the Round 2 submission as the rules require.

PrismCheck AI · Round 1: The Plan · Team A2Z Creators